

Tensor Flow Cheat Sheet

Becoming Human.AI



In May 2017 Google announced the second-generation of the TPU, as well as the availability of the TPUs in Google Compute Engine.[12] The second-generation TPUs deliver up to 180 teraflops of performance, and when organized into clusters of 64 TPUs provide up to 11.5 petaflops.

Info

TensorFlow

TensorFlow™ is an open source software library created by Google for numerical computation and large scale computation. Tensorflow bundles together Machine Learning, Deep learning models and frameworks and makes them useful by way of common metaphor.

Keras

Keras is an open sourced neural networks library, written in Python and is built for fast experimentation via deep neural networks and modular design. It is capable of running on top of TensorFlow, Theano, Microsoft Cognitive Toolkit, or PlaidML.

Skflow

Scikit Flow is a high level interface base on tensorflow which can be used like sklearn. You can build you own model on your own data quickly without rewriting extra code. provides a set of high level model classes that you can use to easily integrate with your existing Scikit-learn pipeline code.

Installation

How to install new package in Python

`pip install <package-name>`

Example: `pip install requests`

How to install tensorflow?

device = `cpu/gpu`

python_version = `cp27/cp34`

`sudo pip install
https://storage.googleapis.com/tensorflow/linux/$device/tensorflow-0.8.0-$python_version-none-linux_x86_64.whl
sudo pip install`

How to install Skflow

`pip install sklearn`

How to install Keras

`pip install keras`

update `~/keras/keras.json` – replace “theano” by “tensorflow”

Helpers

Python helper Important functions

`type(object)`

Get object type

`help(object)`

Get help for object (list of available methods, attributes, signatures and so on)

`dir(object)`

Get list of object attributes (fields, functions)

`str(object)`

Transform an object to string object?

Shows documentations about the object

`globals()`

Return the dictionary containing the current scope's global variables.

`locals()`

Update and return a dictionary containing the current scope's local variables.

`id(object)`

Return the identity of an object. This is guaranteed to be unique among simultaneously existing objects.

`import _builtin_`

`dir(_builtin_)`

Other built-in functions

Tensor Flow

Main classes

`tf.Graph()`

`tf.Operation()`

`tf.Tensor()`

`tf.Session()`

Some useful functions

`tf.get_default_session()`

`tf.get_default_graph()`

`tf.reset_default_graph()`

`ops.reset_default_graph()`

`tf.device("/cpu:0")`

`tf.name_scope(value)`

`tf.convert_to_tensor(value)`

TensorFlow Optimizers

`GradientDescentOptimizer`

`AdadeltaOptimizer`

`AdagradOptimizer`

`MomentumOptimizer`

`AdamOptimizer`

`FtrlOptimizer`

`RMSPropOptimizer`

Reduction

`reduce_sum`

`reduce_prod`

`reduce_min`

`reduce_max`

`reduce_mean`

`reduce_all`

`reduce_any`

`accumulate_n`

Activation functions

`tf.nn?`

`relu`

`relu6`

`elu`

`softplus`

`softsign`

`dropout`

`bias_add`

`sigmoid`

`tanh`

`sigmoid_cross_entropy_with_logits`

`softmax`

`log_softmax`

`softmax_cross_entropy_with_logits`

`sparse_softmax_cross_entropy_with_logits`

`weighted_cross_entropy_with_logits`

etc.

Skflow

Main classes

`TensorFlowClassifier`

`TensorFlowRegressor`

`TensorFlowDNNClassifier`

`TensorFlowDNNRegressor`

`TensorFlowLinearClassifier`

`TensorFlowLinearRegressor`

`TensorFlowRNNClassifier`

`TensorFlowRNNRegressor`

`TensorFlowEstimator`

Each classifier and regressor have following fields
`n_classes=0` (Regressor), `n_classes` are expected to be input (Classifier)

`batch_size=32`,

`steps=200`, // except

`TensorFlowRNNClassifier` - there is 50

`optimizer=Adagrad`,

`learning_rate=0.1`,

Each class has a method fit

`fit(X, y, monitor=None, logdir=None)`

X: matrix or tensor of shape `[n_samples, n_features...]`. Can be iterator that returns arrays of features. The training input samples for fitting the model.

Y: vector or matrix `[n_samples]` or `[n_samples, n_outputs]`. Can be iterator that returns array of targets. The training target values (class labels in classification, real numbers in regression).

monitor: Monitor object to print training progress and invoke early stopping

logdir: the directory to save the log file that can be used for optional visualization.

`predict(X, axis=1, batch_size=None)`

Args:

X: array-like matrix, `[n_samples, n_features...]` or iterator.
axis: Which axis to argmax for classification.

By default axis 1 (next after batch) is used. Use 2 for sequence predictions.

`batch_size`: If test set is too big, use batch size to split it into mini batches. By default the `batch_size` member variable is used.

Returns:

y: array of shape `[n_samples]`. The predicted classes or predicted value.